

Graphical Models



Adrian Barbu

What is a Graphical Model?

Example: Medical Diagnosis

■ Observed data:

- Symptoms
- Medical tests
- Circumstances

Visit to Asia X_1

Smoking X_2

Tuberculosis X_3

Lung Cancer X_4

■ Unknowns:

- Cause (disease)

Tuberculosis or Cancer X_6

Bronchitis X_5

XRay Result X_7

Dyspnea X_8

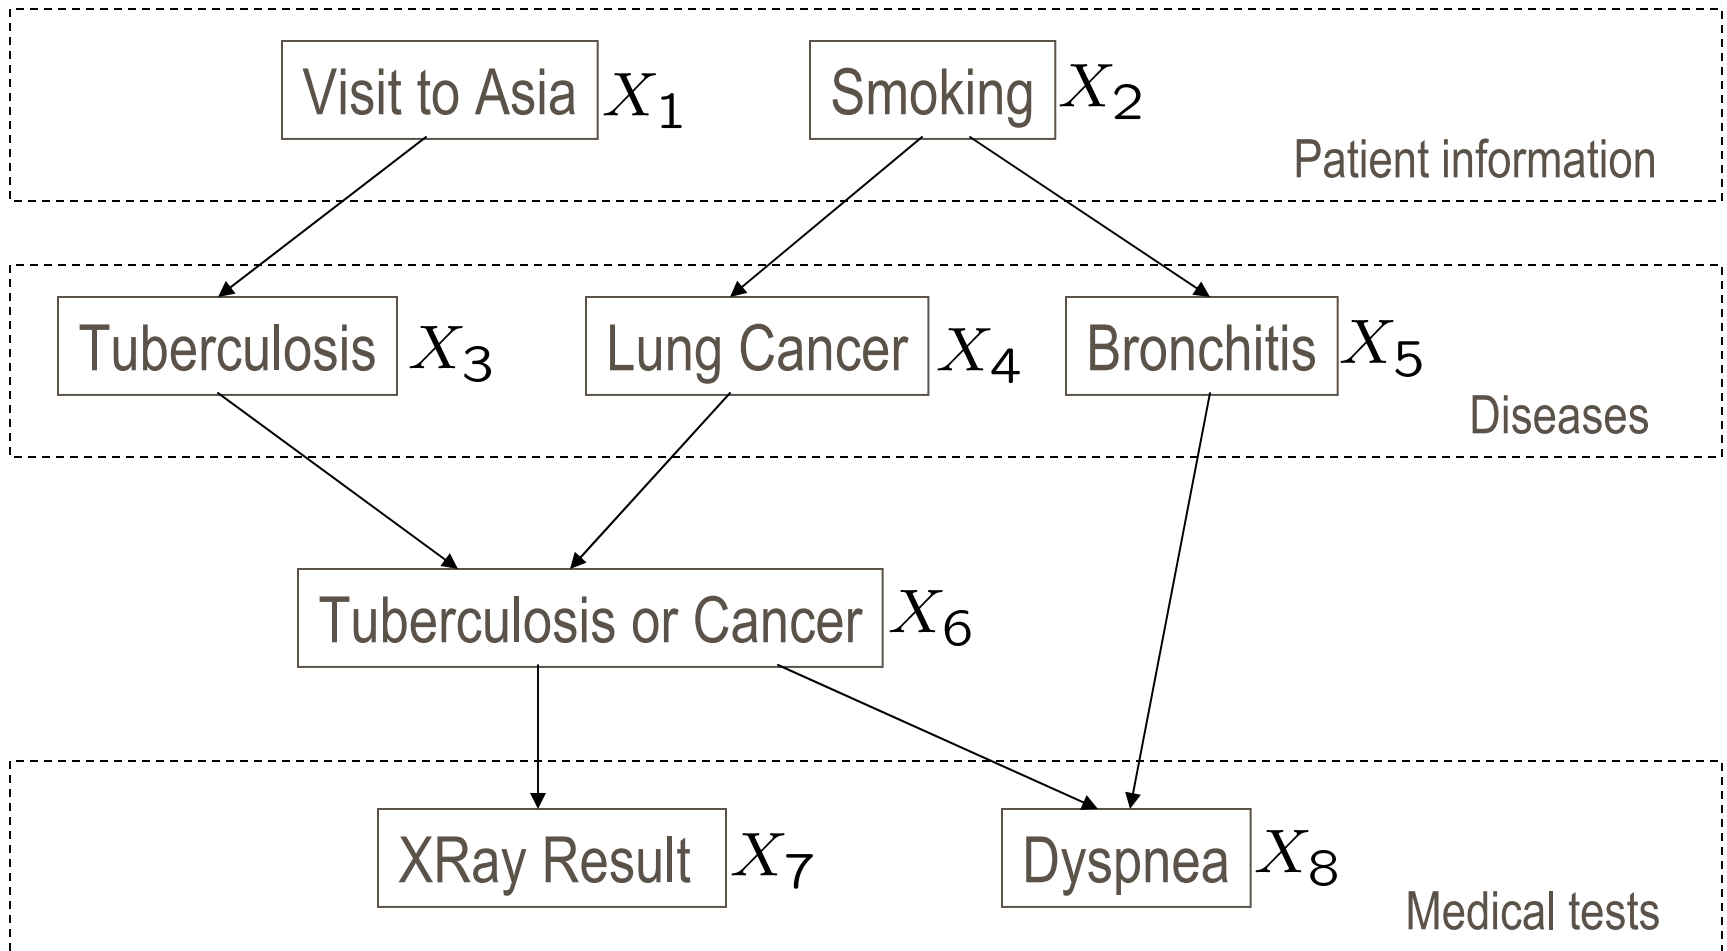
Probabilities

- Want to find most probable disease given the observations
- Can represent whole process by a joint probability

$$P(X_1, X_2, X_3, X_4, X_5, X_6, X_7, X_8)$$

- 2^8 states
- Not all are needed
- No medical insight
- Need too much training data
- Encode domain knowledge into relationships between variables
 - Need less training data
- Inference
 - Find most probable unknown states given observations

Dependency Graph

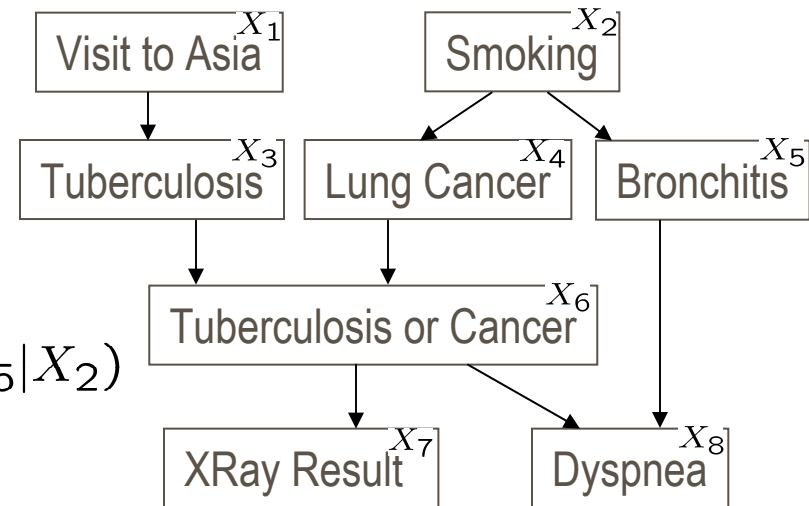


Probabilistic Graphical Models

- Representation of the dependencies using a graph
 - Nodes= random variables
 - Edges= dependencies among variables
 - Edge absence= conditional independence
 - Directed: cause→effect
 - Undirected: bidirectional relationship

- Joint probability factors:

$$\begin{aligned} &P(X_1, X_2, X_3, X_4, X_5, X_6, X_7, X_8) \\ &= P(X_1)P(X_2)P(X_3|X_1)P(X_4|X_2)P(X_5|X_2) \\ &\cdot P(X_6|X_3, X_4)P(X_7|X_6)P(X_8|X_5, X_6) \end{aligned}$$



Advantages of PGM

- Why is this good?
 - Each probability involves fewer variables
 - Need fewer samples to train
 - Not so prone to overfitting
 - Total number of bins is smaller
 - Before were $2^8=256$ bins
 - Now $2+2+4+4+4+8+4+8=36$
 - 8 times fewer bins
 - Easier to find the best solution
 - Explore the graph structure
 - Faster than exhaustive search of all possibilities
 - Incorporation of domain knowledge
 - Avoids impossible solutions

Two types of PGM

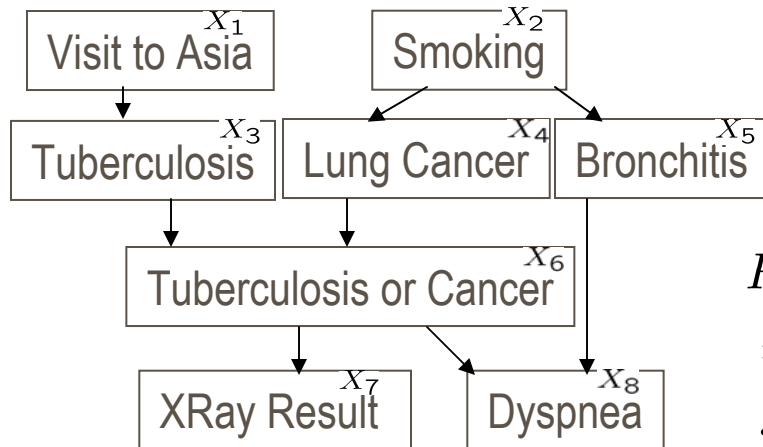
■ Directed edges

- Model causality relationships
- Bayesian Network or Directed Graphical Model
- DAG=Directed Acyclic Graph

■ Undirected edges

- Model symmetric correlations between variables
- Markov Random Field or Undirected Graphical Model

Bayesian Network:



$$\begin{aligned} P(X_1, X_2, X_3, X_4, X_5, X_6, X_7, X_8) \\ = P(X_1)P(X_2)P(X_3|X_1)P(X_4|X_2)P(X_5|X_2) \\ \cdot P(X_6|X_3, X_4)P(X_7|X_6)P(X_8|X_5, X_6) \end{aligned}$$

Factorization Theorem:

- Given a DAG, the most general form of the probability distribution that is **consistent with** the graph factors according to “node given its parents”:

$$P(\mathbf{X}) = \prod_{i=1}^d P(X_i | \mathbf{X}_{\pi_i})$$

- \mathbf{X}_{π_i} is the set of parents of X_i
- d = number of nodes (variables) in the graph

Bayesian Network Properties

■ Markov Blanket

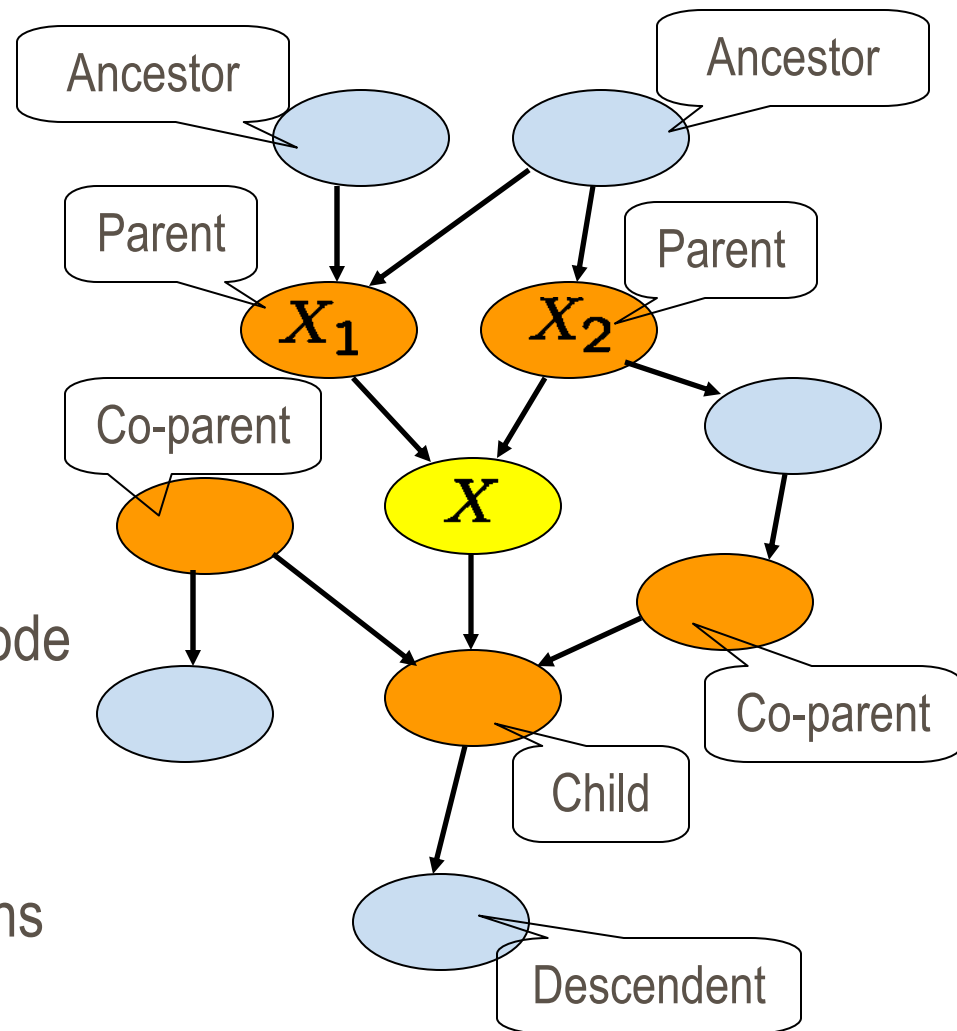
- Parents
- Children
- Co-parents (parents of the children)

■ Conditional Independence

- Any node is conditional independent of any other node not in its Markov Blanket

■ Joint distribution from:


- Local conditional distributions
- DAG structure



Conditional Independence

■ Common parent

- Knowing C decouples A and B
- A and B are conditionally independent given C




■ Chain

- Knowing B decouples A and C
- A and C are conditionally independent given B



■ Common child

- A and B are independent
- Knowing C couples them together
- A or B can explain-away the observation C



Explain Away Example

- Graph nodes:

 - B=battery, F=fuel, G=gauge

- B and F are independent

$$P(B = 1) = 0.9, P(F = 1) = 0.9$$

- Unreliable fuel gauge: $P(G = 1|B = 1, F = 1) = 0.8$


$$P(G = 1|B = 1, F = 0) = 0.2$$

$$P(G = 1|B = 0, F = 1) = 0.2$$


$$P(G = 1|B = 0, F = 0) = 0.1$$

- Suppose we observe $G = 0$

- Want $P(F = 0|G = 0)$



Explain Away Example



- Bayes Rule:

$$P(F = 0|G = 0) = \frac{P(G = 0|F = 0)P(F = 0)}{P(G = 0)}$$

- We compute

$$P(G = 0) = \sum_B \sum_F P(G = 0|B, F)P(B)P(F) = 0.315$$

$$P(G = 0|F = 0) = \sum_B P(G = 0|B, F = 0)p(B) = 0.81$$

- So we get

$$P(F = 0|G = 0) = \frac{0.81 \cdot 0.1}{0.315} = 0.257$$

- If we observe that $B = 0$, we get

$$P(F = 0|G = 0, B = 0) = \frac{P(G = 0|B = 0, F = 0)P(F = 0)}{\sum_F P(G = 0|B = 0, F)P(F)} = 0.111$$

- So $B = 0$ **explained away** the $G = 0$




D-separation

Problem

- Find whether some groups of nodes are conditionally independent

Reason

- Can compute each independent group separately
 - Faster inference
- X is d-separated from Y given Z if one cannot send a “Bayes ball” from X to Y
 - Z can be known (=dark) or not
 - Notation: $\text{dsep}_G(X, Y|Z)$




D-separation and Independence

- Define all independence properties obtained from the d-separation

$$I(G) = \{X \perp\!\!\!\perp Y|Z : \text{dsep}_G(X, Y|Z)\}$$

- Example



$$\begin{aligned} I(G) = \{ & X_1 \perp\!\!\!\perp X_2 \\ & X_2 \perp\!\!\!\perp X_4 \\ & X_2 \perp\!\!\!\perp X_4|X_1 \\ & X_2 \perp\!\!\!\perp X_4|\{X_1, X_3\} \\ & X_3 \perp\!\!\!\perp X_4|X_1 \\ & X_4 \perp\!\!\!\perp \{X_2, X_3\}|X_1 \} \end{aligned}$$


Equivalence Theorem

- A distribution satisfies $I(G)$ if and only if it factors according to the graph G .

Conditional Probability Tables (CPT)

A^0	0.3
A^1	0.7

B^0	0.4
B^1	0.6




	A^0B^0	A^0B^1	A^1B^0	A^1B^1
C^0	0.4	1	0.9	0.8
C^1	0.6	0	0.1	0.2

	C^0	C^1
D^0	0.3	0.5
D^1	0.7	0.5

$$P(A, B, C, D) = P(A)P(B)P(C|A, B)P(D|C)$$


Conditional Probability Density Functions (CPD)

$$A \sim \mathcal{N}(\mu_A, \Sigma_A) \quad B \sim \mathcal{N}(\mu_B, \Sigma_B)$$




$$C \sim \mathcal{N}(A + B, \Sigma_C)$$

$$D \sim \mathcal{N}(\mu_A + C, \Sigma_D)$$



$$P(A, B, C, D) = P(A)P(B)P(C|A, B)P(D|C)$$

Conditionally Independent Observations



- Variables in plate are replicated conditionally independent given model parameters

Example: Gaussian Model

■ Generative model (Naïve Bayes)

■ Bayes Rule


$$P(\mu, \sigma | y_1, \dots, y_N) = \frac{P(y_1, \dots, y_N | \mu, \sigma) P(\mu, \sigma)}{P(y_1, \dots, y_N)}$$

■ i.i.d assumption

■ Model parameters: μ, σ

■ Likelihood = P(Observations|Parameters):

$$P(y_1, \dots, y_N | \mu, \sigma) = \prod_{i=1}^N P(y_i | \mu, \sigma) = \prod_{i=1}^N \frac{1}{\sqrt{2\pi}\sigma} e^{-(y_i - \mu)^2 / 2\sigma^2}$$



Example: Gaussian Model


■ Prior on model parameters

- Assume parameters independent

$$P(\mu, \sigma) = P(\mu|\alpha)P(\sigma|\beta)$$

■ Possible priors $P(\mu|\alpha)$

- Parametric:
 - Gaussian
 - Student t distribution
 - ...
- Non-parametric
 - Histogram
 - Parzen windows



Markov Random Fields

■ Undirected Graph:


- A node is conditionally independent of every other node given its direct neighbors

■ Probability Distribution


- A set of cliques = complete subgraphs
- A set of potential functions on the cliques

■ Pros and cons:

- Easy to compute the probability
- Hard to obtain samples from the probability
- Hard to find max probability configuration



Conditional Independence




G undirected graph

- B *separates* A and C if every path from A to C passes through a node in B $\text{sep}_G(A, C|B)$
- A probability distribution satisfies the *Global Markov Property* if for any disjoint A, B, C , such that B separates A and C then A is independent of C given B :

$$I(G) = \{A \perp\!\!\!\perp C|B : \text{sep}_G(A, C|B)\}$$

Cliques

- Clique:
 - A complete subgraph $G'=(V',E')$ of $G=(V,E)$
 - Complete means fully connected
- Maximal clique
 - A clique such that there is no other clique that includes it
- Example:
 - Max-cliques:
 - $\{A,B,C\}, \{B,C,D\}$
 - Other cliques:
 - $\{A,B\}, \{A,C\}, \{B,C\}$
 - $\{B,D\}, \{C,D\}$
 - $\{A\}, \{B\}, \{C\}, \{D\}$



Markov Random Field Probability

Given:

- An undirected graph G
- Define a set C of cliques of G
- Define a set of potential functions $\psi_c(\mathbf{x}_c)$, $\forall c \in C$
 - Encourage certain configurations

MRF probability:

$$P(\mathbf{x}) = P(x_1, \dots, x_N) = \frac{1}{Z} \prod_{c \in C} \psi_c(\mathbf{x}_c)$$

where Z is the partition function


$$Z = \sum_{\mathbf{x}} \prod_{c \in C} \psi_c(\mathbf{x}_c)$$

Example with Max-cliques

- C=all max-cliques

$$P(x_1, \dots, x_4) = \frac{1}{Z} \psi_1(x_1, x_2, x_3) \psi_2(x_2, x_3, x_4)$$

$$Z = \sum_{x_1, x_2, x_3, x_4} \psi_1(x_1, x_2, x_3) \psi_2(x_2, x_3, x_4)$$




- For discrete x_1, \dots, x_4 , we can represent $P(x_1, \dots, x_4)$ with two 3D tables (histograms) instead of one 4D table.
 - Avoid overfitting

Example with Sub-cliques

- C=all 2D cliques


$$P(x_1, \dots, x_4) = \frac{1}{Z} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \\ \cdot \psi_{23}(x_2, x_3) \psi_{24}(x_2, x_4) \psi_{34}(x_3, x_4)$$



$$Z = \sum_{x_1, x_2, x_3, x_4} \psi_{12}(x_1, x_2) \psi_{13}(x_1, x_3) \psi_{23}(x_2, x_3) \psi_{24}(x_2, x_4) \psi_{34}(x_3, x_4)$$

- For discrete x_1, \dots, x_4 , we can represent $P(x_1, \dots, x_4)$ with five 2D tables instead of one 4D table.
 - Avoid overfitting
 - But it might be a weaker model

Clique Potentials



- Clique potentials are not probabilities

$$\begin{aligned} P(A, B, C) &= P(A)P(B|A)P(C|B) \\ &= P(B)P(A|B)P(C|B) \\ &= P(C)P(B|C)P(A|B) \\ &= P(A, B)P(C|B) \end{aligned}$$

- We cannot have only marginal probabilities, e.g. $P(A)$
- We cannot have only conditional probabilities, e.g. $P(B|A)$
- Probability = all clique potentials + partition function Z

Exponential Form

- Since all cliques are positive, can use exponential form

$$\psi_c(\mathbf{x}_c) = \exp[-\phi_c(\mathbf{x}_c)]$$

- $\phi_c(\mathbf{x}_c)$ is also called potential
- Obtain exponential form of the probability

$$P(\mathbf{x}) = \frac{1}{Z} \exp\left[-\sum_{c \in C} \phi_c(\mathbf{x}_c)\right] = \frac{1}{Z} \exp[-H(\mathbf{x})]$$

- $H(\mathbf{x})$ is called the free energy

$$H(\mathbf{x}) = \sum_{c \in C} \phi_c(\mathbf{x}_c)$$

- Exponential form

- In physics is called Boltzmann distribution
- In statistics, it is called log-linear model

Example: Boltzmann Machine


Boltzmann Machine:

- Binary node values $x_i = \pm 1$
- Fully connected graph G
- Pairwise potentials $\phi_{ij}(x_i, x_j)$
- Probability distribution:

$$\begin{aligned} P(x_1, \dots, x_4) &= \frac{1}{Z} \exp\left[-\sum_{i,j} \phi_{ij}(x_i, x_j)\right] \\ &= \frac{1}{Z} \exp\left[-\sum_{i,j} \theta_{ij} x_i x_j - \sum_i \alpha_i x_i\right] \end{aligned}$$


- Energy:

$$H(\mathbf{x}) = \sum_{i,j} (x_i - \mu_i) \theta_{ij} (x_j - \mu_j) = (\mathbf{x} - \mu)^T \Theta (\mathbf{x} - \mu)$$



Example: Ising/Potts Models

- Nodes on a grid
- Node values (labels):
 - ± 1 for Ising model
 - $1 \dots N > 2$ for Potts model
- Edges only with the 4 neighbors



$$P(\mathbf{x}) = \frac{1}{Z} \exp\left[-\sum_{i,j} \theta_{ij} I(x_i \neq x_j) - \sum_i \alpha_i x_i\right]$$

- E.g.
 - Nodes are pixels
 - $\theta_{ij} = \theta > 0$ encourages nearby pixels to have same label
 - α_i form the “external field” (the data term, the likelihood)

Application: Image Denoising

■ Ising model $P(\mathbf{x}) = \frac{1}{Z} \exp[-\theta \sum_{i,j} x_i x_j - \alpha \sum_i x_i - \eta \sum_i x_i y_i]$

■ Parameters

■ $\theta = 1$


■ $\alpha = 0$

■ $\eta = 2.1$


■ Result
depends on:

■ Model


■ Algorithm




Original image



Noisy image



Restored with ICM



Restored with min-cut

Example: Protein-Protein Interaction Networks

C:\Documents ar	38	126
C:\Shubhada\ne	1681	3053
C:\Shubhada	35	40
C:\Shubhada	49	281
C:\Shubhada	10	18

Cell Cycle Signal Transduction

Carbon metabolism

Hydrogen Transport

Biosynthesis protein

Ribosomal Complex

Energy cox

Protein Metabolism

Vesicle Transport

Cell Cycle Regulation

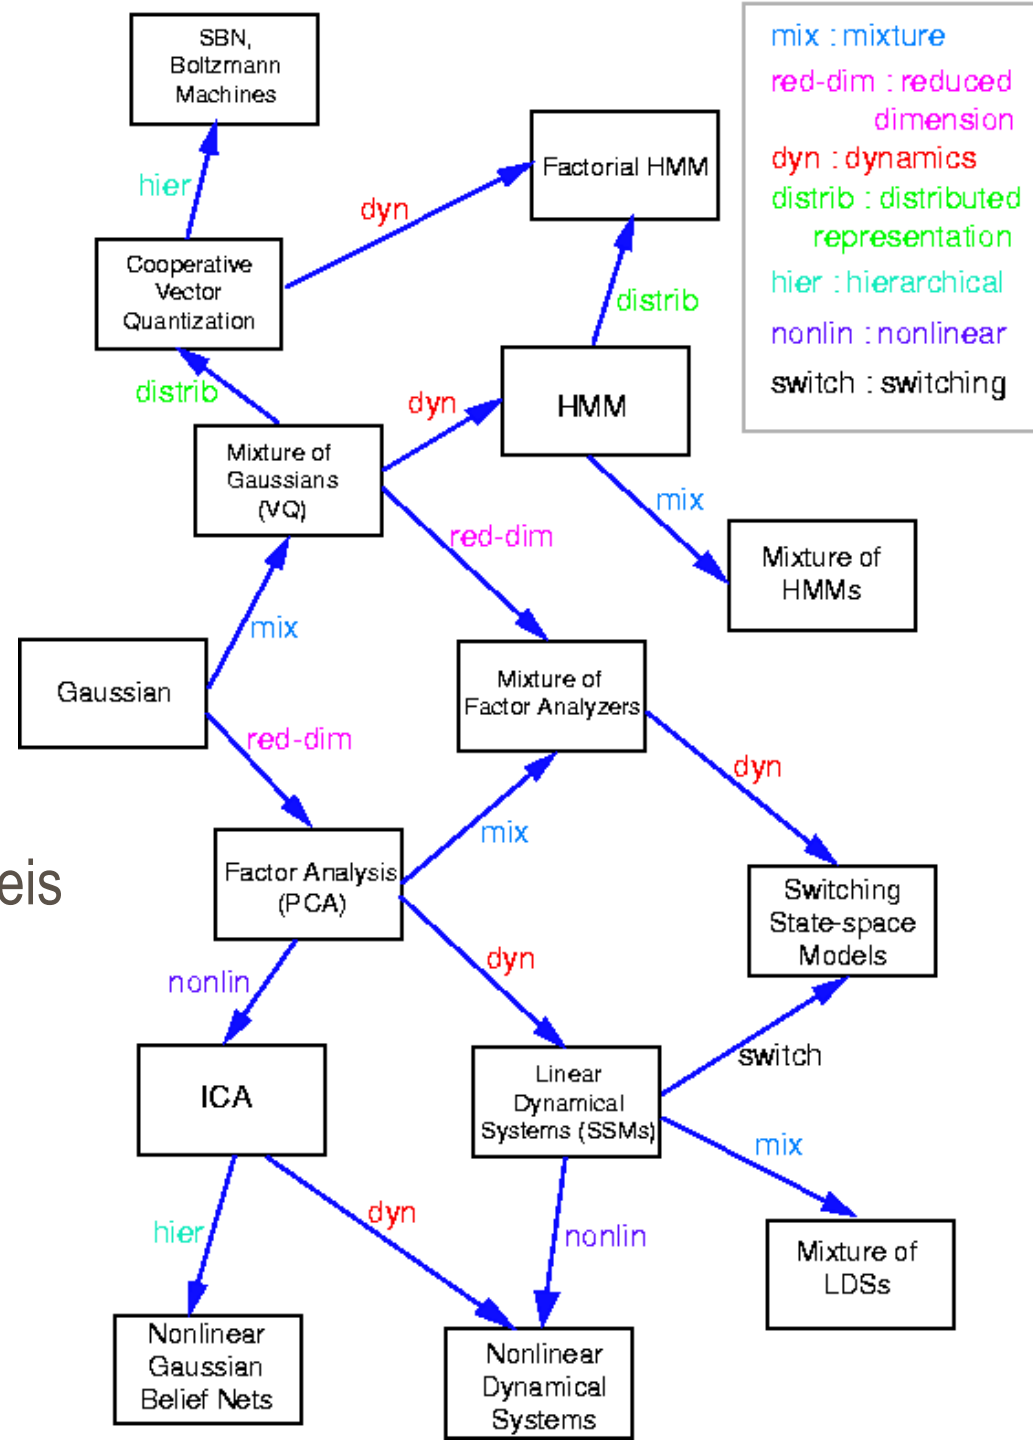
Protanprt_TOM

Ion homeostasis

DB



Genealogy of Graphical Models




Zoubin Ghahramani and Sam Roweis

<http://www.cs.ubc.ca/~murphyk/Bayes/bnintro.html>



Learning Graphical Models

Task:

- Given training data
- Find “best” DAG and CPD



$(A, B, C, D, E) = (00110)$
 $(A, B, C, D, E) = (00010)$
 $(A, B, C, D, E) = (00111)$
 ...
 $(A, B, C, D, E) = (01110)$



	A^0B^0	A^0B^1	A^1B^0	A^1B^1
C^0	0.4	1	0.9	0.8
C^1	0.6	0	0.1	0.2

A^0	0.3
A^1	0.7

B^0	0.4
B^1	0.6

Conditional Probability Tables

Bayesian Learning Approach

- Θ = all unknown parameters
 - Graph structure (preferably not)
 - Probabilities
- Bayes Rule $P(\Theta|y_1, \dots, y_N) \propto P(y_1, \dots, y_N|\Theta)P(\Theta)$
- Find the mean Θ
$$\Theta = \int \Theta P(\Theta|y_1, \dots, y_N) d\Theta$$
- Sampling Θ
 - Use MCMC
 - Tricky if graph structure is in Θ
 - Sometimes use uniform $P(\Theta)$

Inference

Graphical Models

- Compact Representation of a probability distribution P

Inference:

- Answer queries about P .
- Examples:
 - Is node X independent on node Y given nodes Z, W ?
 - Probability of $X=\text{true}$ if $(Y=\text{false} \text{ and } Z=\text{true})=?$
 - What is the joint distribution of (X, Y) if $Z=\text{false}$?
 - What is the likelihood when assigning values to all variables?
 - What is the most likely assignment of values to all or a subset the nodes, knowing other nodes?

Query 1: Likelihood

- Some variables have been observed
 - Say the variables are $E = \{x_{k+1}, \dots, X_n\}$
 - They form the **evidence**
 - Are assigned some value vector e
- Compute probability of the evidence

$$P(e) = \sum_{x_1} \dots \sum_{x_k} P(x_1, \dots, x_k, e)$$

- Aka likelihood of e
- Need to integrate all other variables
- Need to know the partition function (for MRF)

Query 2: Conditional Probability

- Conditional probability distribution of the remaining variables given the evidence

$$P(X|e) = \frac{P(X, e)}{P(e)} = \frac{P(X, e)}{\sum_x P(X = x, e)}$$

the *a posteriori belief* in X , given evidence e

- If we query a subset Y of the variables $X=\{Y,Z\}$, we integrate out (don't care about) Z :

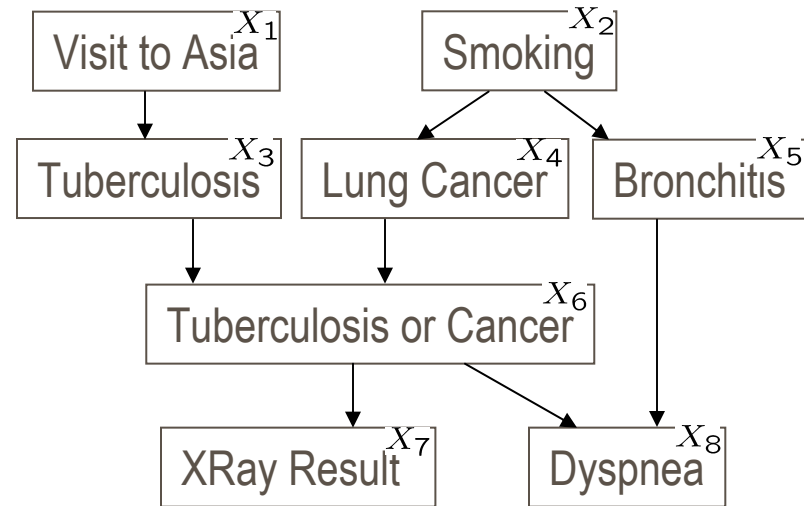
$$P(Y|e) = \sum_z P(Y, Z = z|e)$$

- integrate out = **marginalization**

Applications of the A-posteriori Belief

■ Prediction:

- Probability of an outcome given the starting condition
- The query node is a descendent of the evidence
- E.g.: Probability of lung cancer, if smoking



■ Diagnosis:

- Probability of disease/fault given symptoms
- The query node an ancestor of the evidence
- E.g.: Probability of cancer given X-ray result and dyspnea

■ Direction of information flow \neq direction of edges in GM

- Inference combines evidence from all parts of the graph

Query 3: Most Probable Assignment

- Find most probable values for variables Y
- Given evidence e
- Ignore the rest Z of variables

$$MPA(Y|e) = \arg \max_y P(y|e) = \arg \max_y \sum_z P(y, z|e)$$

- The **maximum a posteriori** assignment for Y given e

Applications of MPA

Classification

- Find most likely label, given the evidence

Explanation

- Find the most likely scenario, given the evidence

Observations:

- The MPA of a variable depends on its "context"- the set of variables being jointly queried
- Example:
 - MPA of $X = 1$
 - MPA of $(X, Y) = (0, 0)$

X	Y	P(X,Y)
0	0	0.4
0	1	0.05
1	0	0.25
1	1	0.3

Inference Complexity

Theorem:

- For a general GM, computing $P(X = x \mid e)$ is NP-hard.
- We cannot find a general procedure that works efficiently for arbitrary GMs


But

- Approximate (suboptimal) solutions always exist
- For particular families of GMs, there are provably efficient exact procedures

Inference Algorithms

- Exact algorithms
 - The elimination algorithm
 - The junction tree algorithms
 - Efficient for some graphs, otherwise exponential
- Approximate inference techniques
 - Stochastic simulation / sampling methods
 - Markov chain Monte Carlo methods
 - Variational algorithms
 - Belief Propagation
 - Tradeoff speed/accuracy

Variable Elimination on Chains




- Say E is observed $E=e$
- Query: Compute $P(e)$

$$P(e) = \sum_a \sum_b \sum_c \sum_d P(a, b, c, d, e)$$

- Exponential number of terms
- Using the chain structure we have

$$\begin{aligned} P(e) &= \sum_d \sum_c \sum_b \sum_a P(a)P(b|a)P(c|b)P(d|c)P(e|d) \\ &= \sum_d \sum_c \sum_b P(c|b)P(d|c)P(e|d) \sum_a P(a)P(b|a) \end{aligned}$$

Variable Elimination on Chains



- Actually $\sum_a P(a)P(b|a) = P(b)$
- Compute and memorize all $P(b)$


■ We eliminated A

■ We now have

$$\begin{aligned} P(e) &= \sum_d \sum_c \sum_b P(c|b)P(d|c)P(e|d)P(b) \\ &= \sum_d \sum_c P(d|c)P(e|d) \sum_b P(c|b)P(b) \end{aligned}$$

- Compute and memorize all $P(c) = \sum_b P(b)P(c|b)$
- We eliminated B

Variable Elimination on Chains



- Repeat the same trick for C and D

- Compute and memorize all $P(c)$
- Then compute and memorize all $P(d)$

- Now we get what we want

$$P(e) = \sum_d P(e|d)P(d)$$

- Similar to dynamic programming

- Save computation by memorization

- Time complexity

- Each step costs $O(|Val(X_i)| \cdot |Val(X_{i+1})|)$ operations: $O(kn^2)$
- Brute force= $O(n^k)$ ($n=|Val(X_i)|$, k = chain length)

Variable Elimination on General DGMs

Idea:

- Write the full probability $P(x_1, \dots, x_n) = \prod P(x_i | \text{parents}(X_i))$
- Integrate out the variables that are not in the query and are not observed

$$P(x_1, e) = \sum_{x_k} \dots \sum_{x_2} \prod_i P(x_i | \text{parents}(X_i))$$

- Choose a good elimination order
- Iterate
 - Move all irrelevant terms outside of innermost sum
 - Compute innermost sum and memorize its values
 - Insert the new term into the product
- Obtain the final result

$$P(x_1 | e) = \frac{P(x_1, e)}{P(e)}$$

Example

- Query: $P(b|h)$

- Full probability



$$P(a, b, c, d, e, f, g, h) = P(a)P(b)P(c|a) \\ \cdot P(d|b)P(f|b)P(e|c, d)P(g|e)P(h|e, f)$$

- Need to eliminate A,C,D,E,F,G,H
- Start with H, memorize the marginal

$$m_h(e, f) = P(h_0|e, f)$$

- Obtain

$$P(a, b, c, d, e, f, g, h_0) = P(a)P(b)P(c|a) \\ \cdot P(d|b)P(f|b)P(e|c, d)P(g|e)m_h(e, f)$$



Example

- Memorize $m_g(e) = \sum_g P(g|e)$

- Eliminate **G**

$$P(a, b, c, d, e, f, h_0) = \sum_g P(a)P(b)P(c|a)$$

$$\cdot P(d|b)P(f|b)P(e|c, d)P(g|e)m_h(e, f)$$

$$= P(a)P(b)P(c|a)P(d|b)P(f|b)$$



$$\cdot P(e|c, d)m_g(e)m_h(e, f)$$

- Memorize $m_a(c) = \sum_a P(c|a)P(a)$

- Eliminate **A**

$$P(b, c, d, e, f, h_0) = P(b)P(d|b)P(f|b)$$

$$\cdot P(e|c, d)m_a(c)m_g(e)m_h(e, f)$$



Example

- Memorize $m_f(b, e) = \sum_f P(f|b)m_h(e, f)$



- Eliminate **F**

$$P(b, c, d, e, h_0) = P(b)P(d|b)P(e|c, d) \\ \cdot m_a(c)m_g(e)m_f(b, e)$$

- Memorize $m_c(d, e) = \sum_c P(e|c, d)m_a(c)$

- Eliminate **C**

$$P(b, d, e, h_0) = P(b)P(d|b)m_c(d, e)m_g(e)m_f(b, e)$$



Example

- Memorize $m_e(b, d) = \sum_e m_c(d, e)m_g(e)m_f(b, e)$
- Eliminate **E**



$$P(b, d, h_0) = P(b)P(d|b)m_e(b, d)$$

- Memorize $m_d(b) = \sum_d P(d|b)m_e(b, d)$
- Eliminate **D**

$$P(b, h_0) = P(b)m_d(b)$$

- Compute $P(h_0) = \sum_b P(b, h_0)$

- Obtain the final result $P(b|h_0) = \frac{P(b, h_0)}{P(h_0)}$




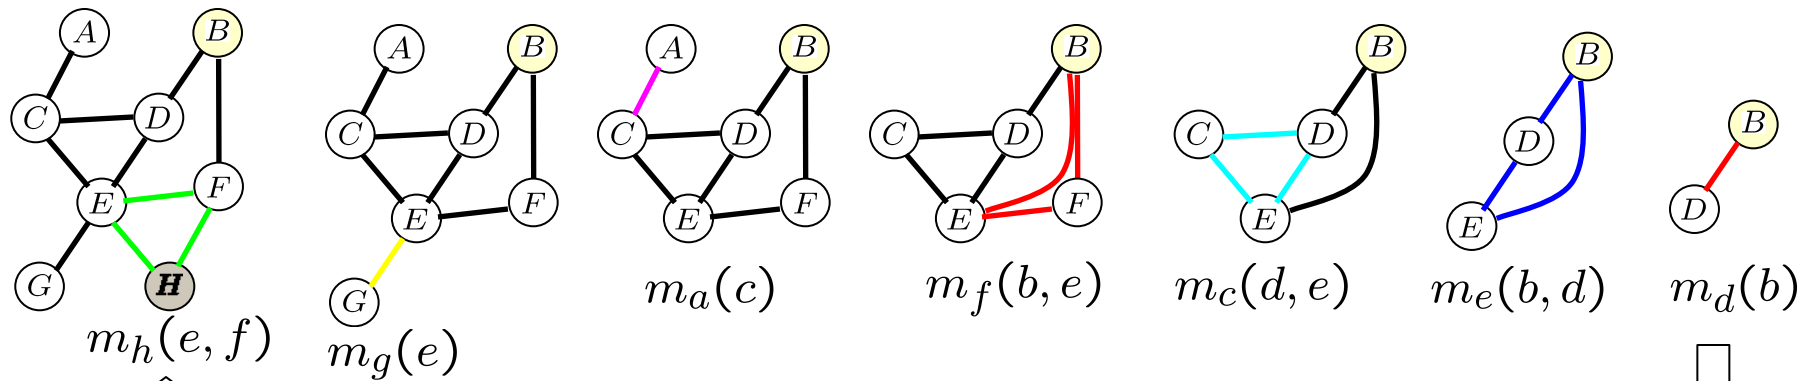
Time Complexity of Variable Elimination

- Look at one elimination step, when we memorize


$$m_x(x_1, \dots, x_k) = \sum_x \prod_{i=1}^k m_i(x, x_{c_i})$$

- We compute this for all values of (x_1, \dots, x_k)
- Memory requirement $\prod_{i=1}^k |Val(x_i)|$
- For each entry, we need $k|Val(x)|$ multiplications and $|Val(x)|$ additions
- Totally
 - $k|Val(x)| \prod_{i=1}^k |Val(x_i)|$ multiplications
 - $|Val(x)| \prod_{i=1}^k |Val(x_i)|$ additions
- Exponential in number of variables in the intermediate factor
 - Prefer elimination orders with few of variables in each factor

Variable Elimination




Moralization=convert to undirected graph by connecting to the Markov Blanket




- There are many elimination orders
- Preferable one with small number of variables in each clique (memorized marginal)
- Finding the optimum ordering is NP hard
 - Heuristics can be used
- Also works for undirected GMs

The Junction-Tree

- Works on undirected GM
 - For directed GM, convert them to undirected by moralization
- Build the junction tree:
 1. Choose an ordering of the nodes and use Node Elimination to obtain a set of elimination cliques.
 2. Build a complete cluster graph over the elimination cliques.
 3. Weight each edge $\{U, V\}$ by $|U \cap V|$ and compute a maximum-weight spanning tree.



Properties of the Junction-Tree



- **Singly connected:** there is exactly one path between each pair of clusters.
- **Covering:** for each clique C of G there is some cluster node N such that $C \subset N$.
- **Running intersection:** for each pair of clusters B and C containing i , each cluster on the unique path between B and C also contains i .
- Different junction trees are obtained using
 - Different elimination orders
 - Different maximum-weight spanning trees.

Decomposable Probabilities

Definition: A factorized probability

$$P(\mathbf{x}) = P(x_1, \dots, x_N) = \frac{1}{Z} \prod_{c \in C} \psi_c(\mathbf{x}_c)$$

is decomposable if there is a junction tree with cluster set C .

- To convert a factorized p to a decomposable probability:
 1. Build a junction tree T for the GM of P .
 2. Initialize the potentials $\psi_c = 1$ for each cluster $c \in T$.
 3. Multiply each potential ψ of p into the cluster potential ψ_c of one cluster that covers its variables.

Note: This is possible only because of the **covering** property.

The Junction Tree Algorithm


■ Input:

- A decomposable probability
- The associated junction tree T .

■ Output: marginal densities of the cliques

Algorithm:

- Each cluster $c \in T$ knows only its local potential ψ_c and its neighbor clusters.
- Each cluster sends one message (potential function) to each neighbor.
- By combining its local potential with the messages it receives, each cluster is able to compute the marginal density of its variables.
- There are different variants of the algorithm based on what messages are passed




Message Passing

Message passing protocol:

- Cluster B sends a message to a neighbor C only after it has received messages from all neighbors except C.

One version:

- Choose one cluster R to be the root so the junction tree is directed.
- Execute Collect(R) and then Distribute(R)
- Collect(C): For each child B of C,
 1. Recursively call Collect(B)
 2. Pass a message $B \rightarrow C$.
- Distribute(C): For each child B of C
 1. Pass a message $C \rightarrow B$
 2. Recursively call Distribute(B).



The Shafer–Shenoy Algorithm

- The message from B to C is

$$\mu_{BC}(x_{B \cap C}) = \sum_{x_{B-C}} \psi_B(x) \prod_{A \in \partial B - \{C\}} \mu_{AB}(x_A)$$

- cluster B computes the product of its local potential ψ_B and the messages from all clusters except C
- marginalizes out all variables that are not in C
- sends the result to C.


Note: μ_{BC} is well-defined because of the tree structure

- The cluster belief at C is

$$\beta_C(x_C) = \psi_C(x_C) \prod_{B \in \partial C} \mu_{BC}(x_{B \cap C})$$

- After all messages have been passed:

$$\beta_C(x_C) \propto P(x_C)$$



The Hugin Algorithm

- Maintain potential functions for each node C and edge E of the Junction Tree. Initialization:

$$\phi_C(x_C) = \psi_C(x_C),$$

$$\phi_E(x_E) = 1$$

- To pass a message from B to C over edge E , update


$$\phi_E^o(x_E) = \phi_E(x_E),$$

$$\phi_E(x_E) = \sum_{x_{B-E}} \phi_B(x_E, x_{B-E}),$$

$$\phi_C(x_C) = \phi_C(x_C) \frac{\phi_E(x_E)}{\phi_E^o(x_E)}$$

- After all messages have been passed:

$$\phi_C(x_C) \propto P(x_C)$$



Complexity of Junction Tree Algorithms

Junction tree algorithms memorize, multiply, and marginalize potentials:

	Tabular	Gaussian
■ storing ψ_C	$O(k^{ C })$	$O(C ^2)$
■ storing ϕ_E	$O(k^{ E })$	$O(E ^2)$
■ updating ϕ_E from B to C	$O(k^{ B-E })$	$O(B-E ^3 B ^2)$
■ number of messages \sim number of clusters	$O(V)$.	

Thus:

- The time and space complexity is dominated by the size of the largest cluster in the junction tree, named the **width** of the junction tree:
 - When using tables: complexity = exponential in the width.
 - When using Gaussians: complexity = cubic in the width.

Conclusion: Junction Tree

A generic exact inference algorithm for any GM

Algorithm

- Construct junction tree: a special clique tree
- Propagate probabilities - a message-passing protocol

Output:

- Marginal probabilities of all cliques
- Solves all queries in a single run

Complexity:

- exponential in the size of the maximal clique
- a good elimination order often leads to small maximal clique

Many well-known algorithms are special cases of Junction Trees

- Forward-backward, Kalman filter, Sum-Product ...

Inference Algorithms

■ Exact inference:

- The elimination algorithm
 - Obtains one marginal probability
 - Fast on chains and trees
 - Otherwise NP hard
- The junction tree algorithms
 - Obtains all marginal probabilities
 - Still NP hard in general

■ Approximate inference

- Stochastic simulation / sampling
- Markov chain Monte Carlo
- Variational algorithms

Monte Carlo Algorithms

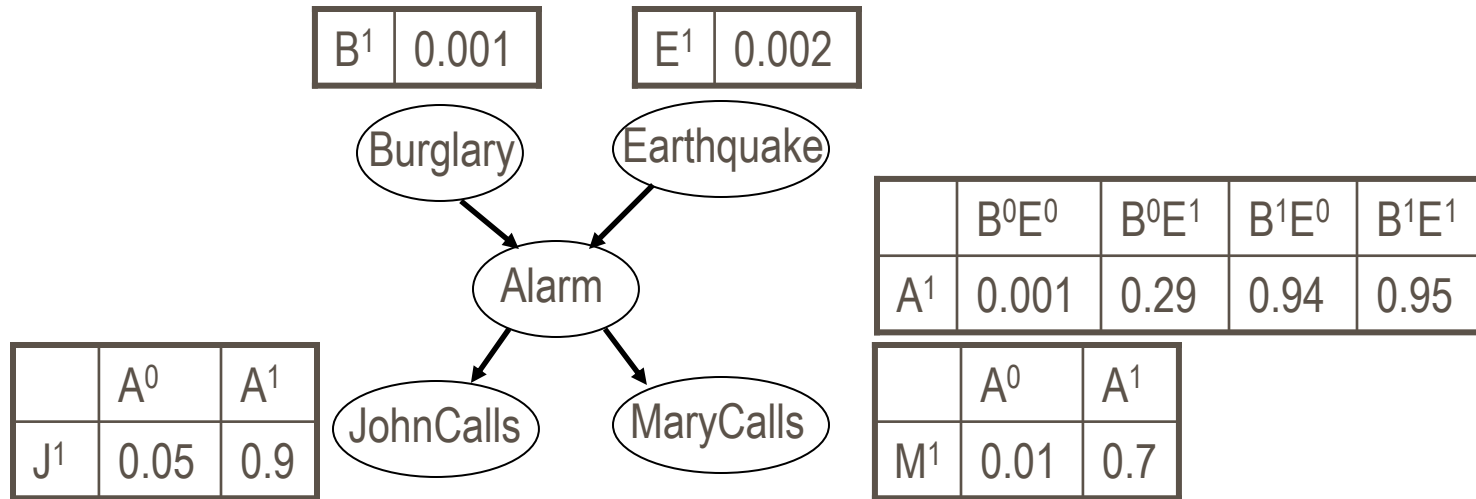
Overview

- Draw random samples from the desired distribution
 - A stochastic representation of a complex distribution
- Marginals and other expectations can be approximated by **sample-based averages**
$$E[f(x)] \approx \frac{1}{N} \sum_{i=1}^N f(x_i)$$
- *Asymptotically* exact and easy to apply to arbitrary models

Challenges:

- Sampling from a given distribution
 - Many distributions are hard to sample (e.g. MRF)
- How to make better use of the samples
 - Not all sample are equally useful
- How to know when to stop sampling

Example: Naïve Sampling

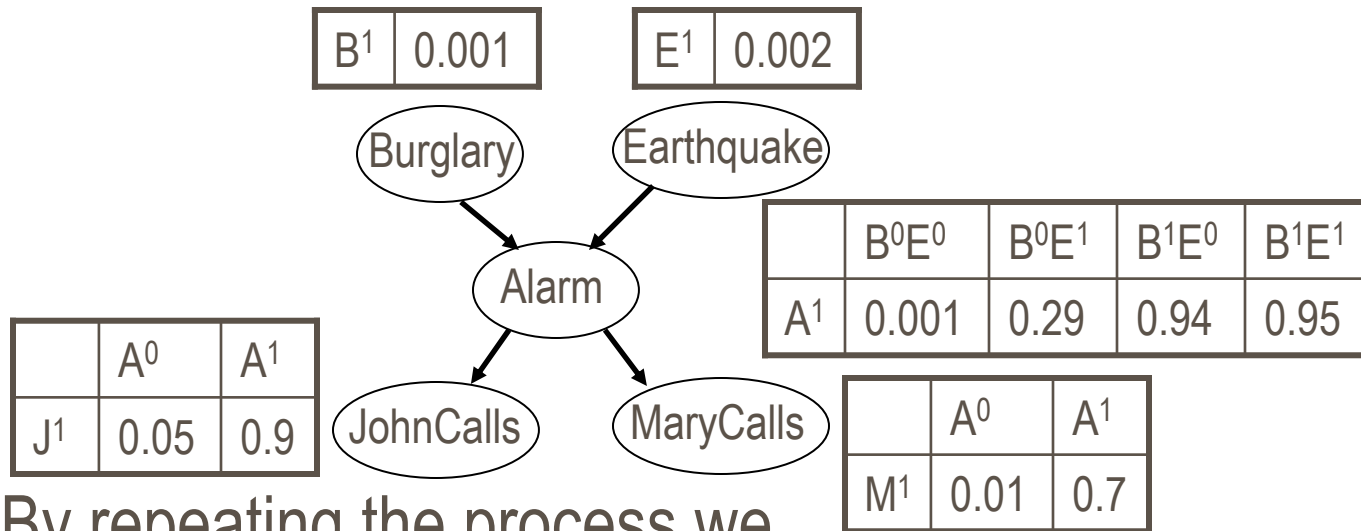


- Sampling: Construct samples according to probabilities given in a Bayesian Network .

Alarm example: (Choose the right sampling sequence)

- 1) Sample from $P(B) = \langle 0.001, 0.999 \rangle$, say we get B^0 .
- 2) Sample E, say we get E^0
- 3) Sample $P(A|B^0, E^0) = \langle 0.001, 0.999 \rangle$, say we get A^0 .

Example: Naïve Sampling



- By repeating the process we get many samples:

- E⁰B⁰A⁰M⁰J⁰
- E⁰B⁰A⁰M⁰J⁰
- E⁰B⁰A⁰M⁰J¹
- E⁰B⁰A⁰M⁰J⁰
- E¹B⁰A¹M¹J¹
- E⁰B⁰A⁰M⁰J⁰

Frequency counting: From the samples we get $P(J|A^0) = P(J, A^0)/P(A^0) = \langle 1/5, 4/5 \rangle$.
Problems:

- What if we want $P(J|A^1)$? We have only one sample
 $P(J|A^1) = P(J, A^1)/P(A^1) = \langle 0, 1 \rangle$.
- What about $P(J|B^1)$? No samples available

Monte Carlo Methods

■ Direct Sampling

- Obtain samples from the GM directly.
- Very difficult to populate a high-dimensional state space

■ Rejection Sampling

- Create samples like direct sampling,
- Only count samples which are consistent with the given evidence.


■ Markov chain Monte Carlo (MCMC)

Markov chain Monte Carlo (MCMC) algorithms

MCMC:


- Iterative stochastic algorithms
- Sample a probability $p(X)$ defined for $X \in \Omega$
- Jump in a new state Y depending on the current state X with probability $K(X, Y)$
- Reversibility:

$$p(X)K(X, Y) = p(Y)K(Y, X), \quad \forall X, Y \in \Omega.$$



Gibbs sampler:

- MCMC algorithm
- Changes one variable X_k at a time by sampling from the marginal probability $p(X_k | X_1, \dots, X_{k-1}, X_{k+1}, \dots, X_n)$

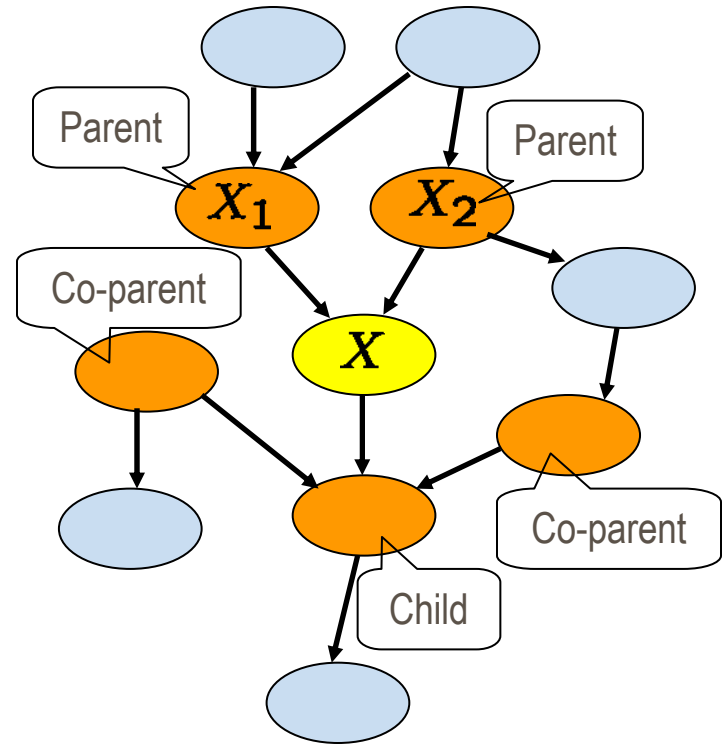


MCMC

Markov-Blanket:

- A node is independent from other nodes, given its parents, children and children's parents.

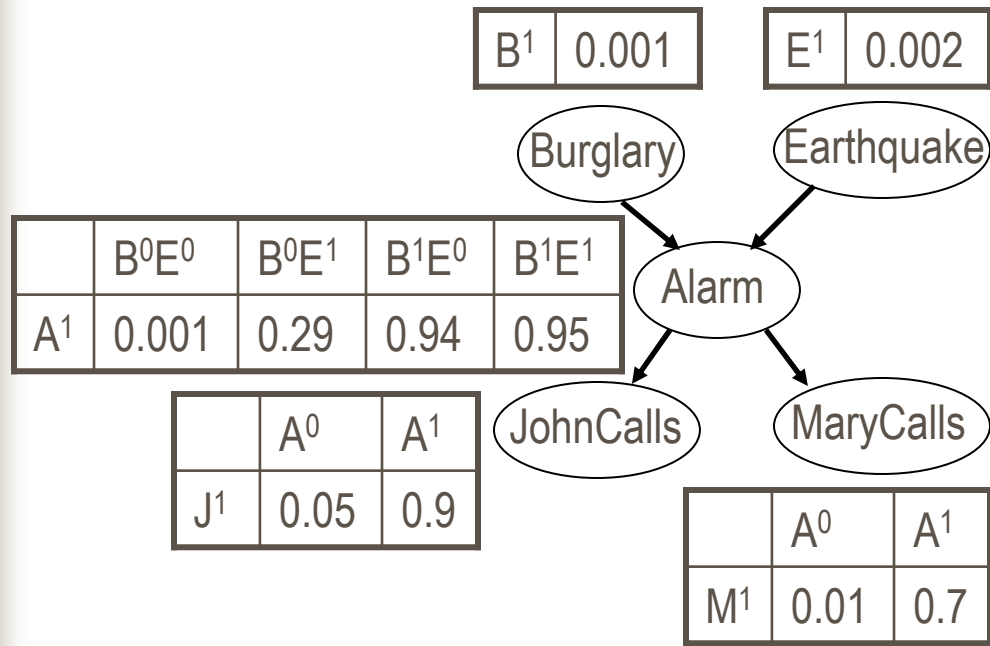
$$p(X_i|X_{-i}) = p(X_i|MB(X_i))$$



Gibbs sampling

- Create a random sample.
- At each step, choose one variable X_i and sample it from $p(X_i|MB(X_i))$ based on the values of the other variables.
- E.g. $MB(A)=\{B, E, J, M\}$, $MB(E)=\{A, B\}$

MCMC Example



Want to calculate $P(J|B^1, M^1)$

- **Evidence** is B^1, M^1 (fixed),

- **Variables** are A, E, J .

- Start from $B^1 E^0 A^1 M^1 J^1$

- Randomly choose next variable (say A)

- Sample A from

$$p(A|MB(A)) = p(A|B^1, E^0, M^1, J^1)$$

say it is 0.

- New state $B^1 E^0 A^0 M^1 J^1$

- Choose next random variable (say E), sample

$$E \sim p(E|MB(E)) = p(E|B^1, A^0)$$

...

- Obtain $P(J|B^1, M^1)$ in two ways:
 - Frequency of $J=1$ during sampling
 - Average of $P(J|A)$ during sampling (faster convergence)

Complexity of MCMC for GM

- Tradeoff speed-accuracy
- It will never reach the true probability, but
 - It will approximate it better and better
 - In an infinite amount of time, it will give the exact solution
- For large and complex graphs, it is much faster than exact inference, since only samples conditionally on the Markov blanket.

References

- M. I. Jordan. Graphical models. *Statistical Science* (Special Issue on Bayesian Statistics), 19, 140-155, 2004.
- <http://www.cs.ubc.ca/~murphyk/Bayes/bnintro.html>
- <http://ai.stanford.edu/~paskin/gm-short-course/lec3.pdf>